GENERAL ASSEMBLY

# DATA SCIENCE

## SYD DAT 6

## Week 6 – Cloud computing, Big Data & Spark
## Wednesday 16th November

1. Cloud Computing
2. Data Stores and Computation
3. Big Data
4. Spark
5. Lab
6. Real World Problem
7. Review

DATA SCIENCE PART TIME COURSE

# CLOUD


APP
OS
APP
OS
APP
OS
APP
OS
APP
OS
APP
OS
Virtualization Platform


amazon
web services


Microsoft
Azure


Google Cloud Platform

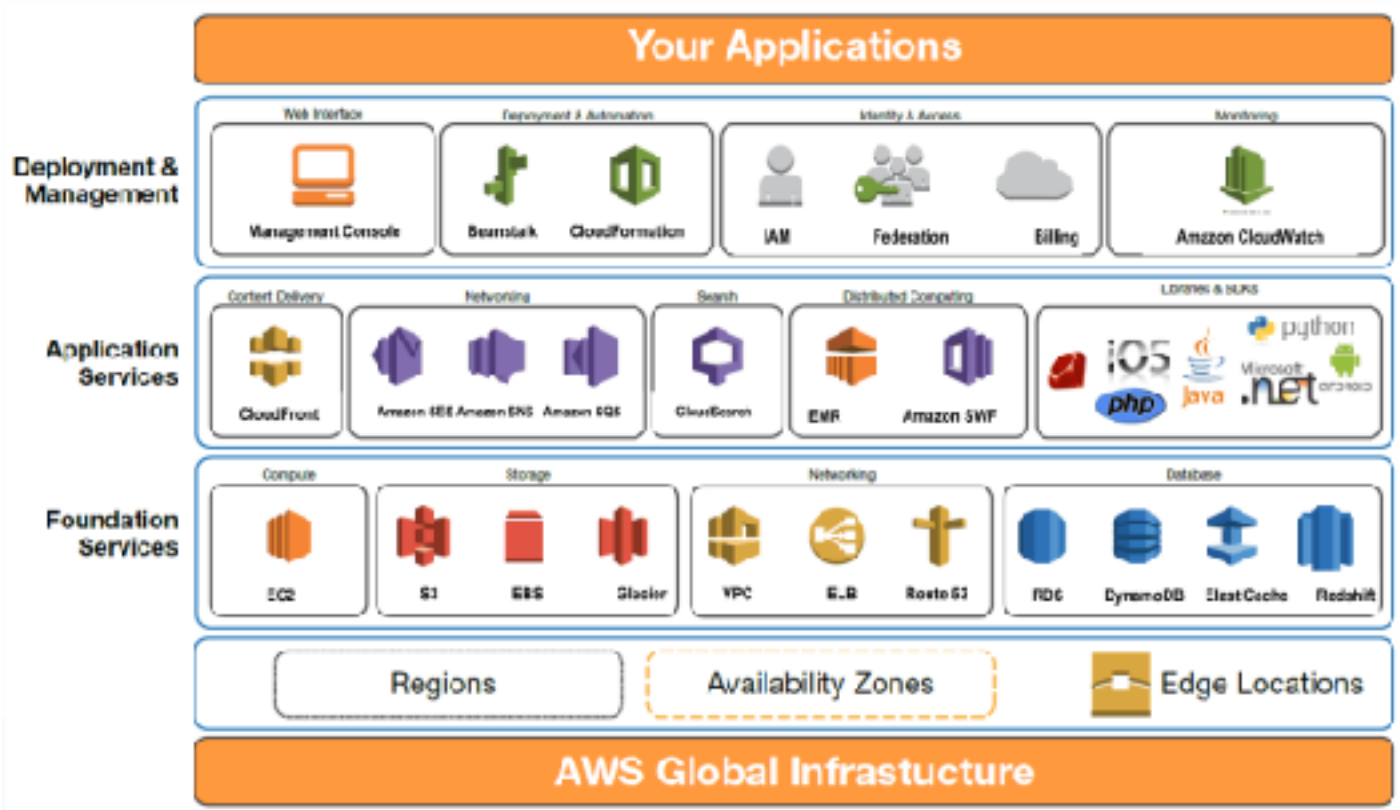
Your Applications
Deployment & Management
Web Interface
Management Console
Deployment & Automation
Beanstalk
CloudFormation
Identity & Access
IAM
Federation
Billing
Monitoring
Amazon CloudWatch
Application Services
Content Delivery
CloudFront
Networking
Amazon SES
Amazon SQS
Search
CloudSearch
Distributed Computing
EMR
Amazon SWF
Libraries & SDKs
ios
python
java
php
.net
Foundation Services
Compute
EC2
Storage
S3
EBS
Glacier
Networking
VPC
ELB
Route 53
Database
RDS
DynamoDB
ElastiCache
Redshift
Regions
Availability Zones
Edge Locations
AWS Global Infrastucture

## https://aws.amazon.com/ec2/pricing/on-demand/

Linux | RHEL | SLES | Windows | Windows with SQL Standard | Windows with SQL Web | Windows with SQL Enterprise

Region: US East (N. Virginia)

| | vCPU | ECU | Memory (GiB) | Instance Storage (GB) | Linux/UNIX Usage |
|---|---|---|---|---|---|
| **General Purpose - Current Generation** | | | | | |
| t2.nano | 1 | Variable | 0.5 | EBS Only | $0.0065 per Hour |
| t2.micro | 1 | Variable | 1 | EBS Only | $0.013 per Hour |
| t2.small | 1 | Variable | 2 | EBS Only | $0.026 per Hour |
| t2.medium | 2 | Variable | 4 | EBS Only | $0.052 per Hour |
| t2.large | 2 | Variable | 8 | EBS Only | $0.104 per Hour |

Docker containers wrap up a piece of software in a complete filesystem that contains everything it needs to run: code, runtime, system tools, system libraries – anything you can install on a server. This guarantees that it will always run the same, regardless of the environment it is running in.

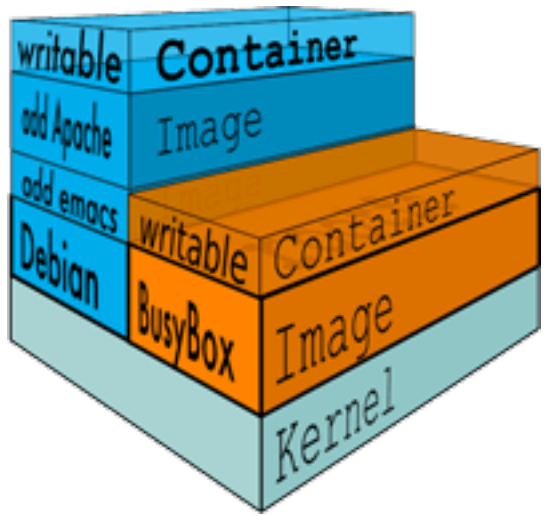


- Lightweight
- Open
- Secure

docker

DATA SCIENCE PART TIME COURSE

# DATA STORES AND COMPUTATION


Microsoft
SQL Server
ORACLE
PostgreSQL
MySQL
SYBASE
IBM
DB2

# MPP RDBMS – Massively Parallel Processing



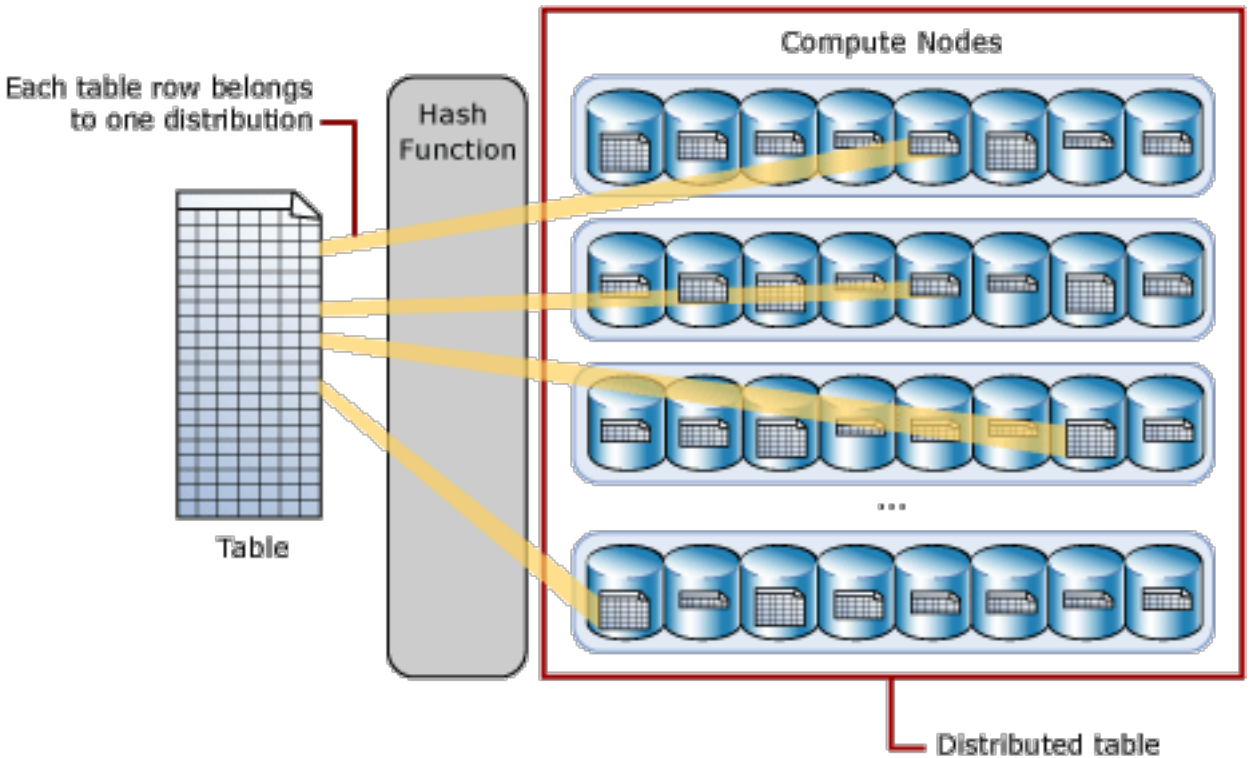


TERADATA.

Greenplum

VERTICA
An HP Company

AMAZON REDSHIFT

## Sources

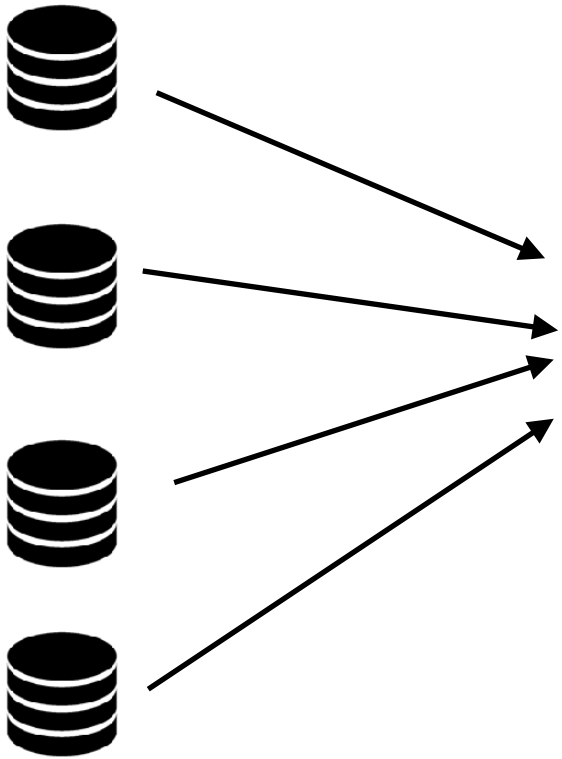

## Structured, relational data

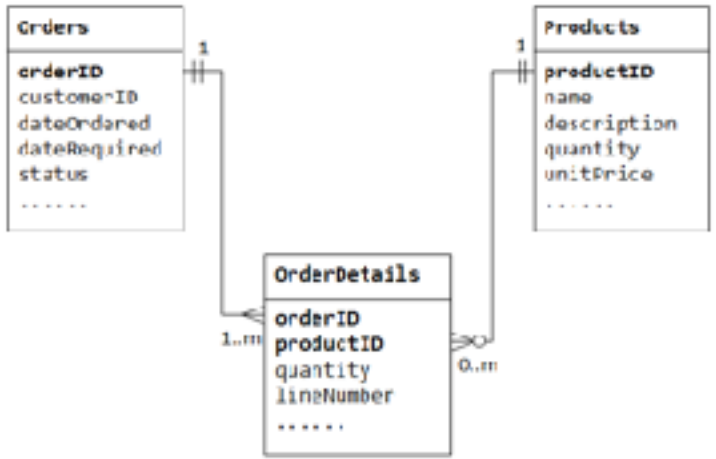

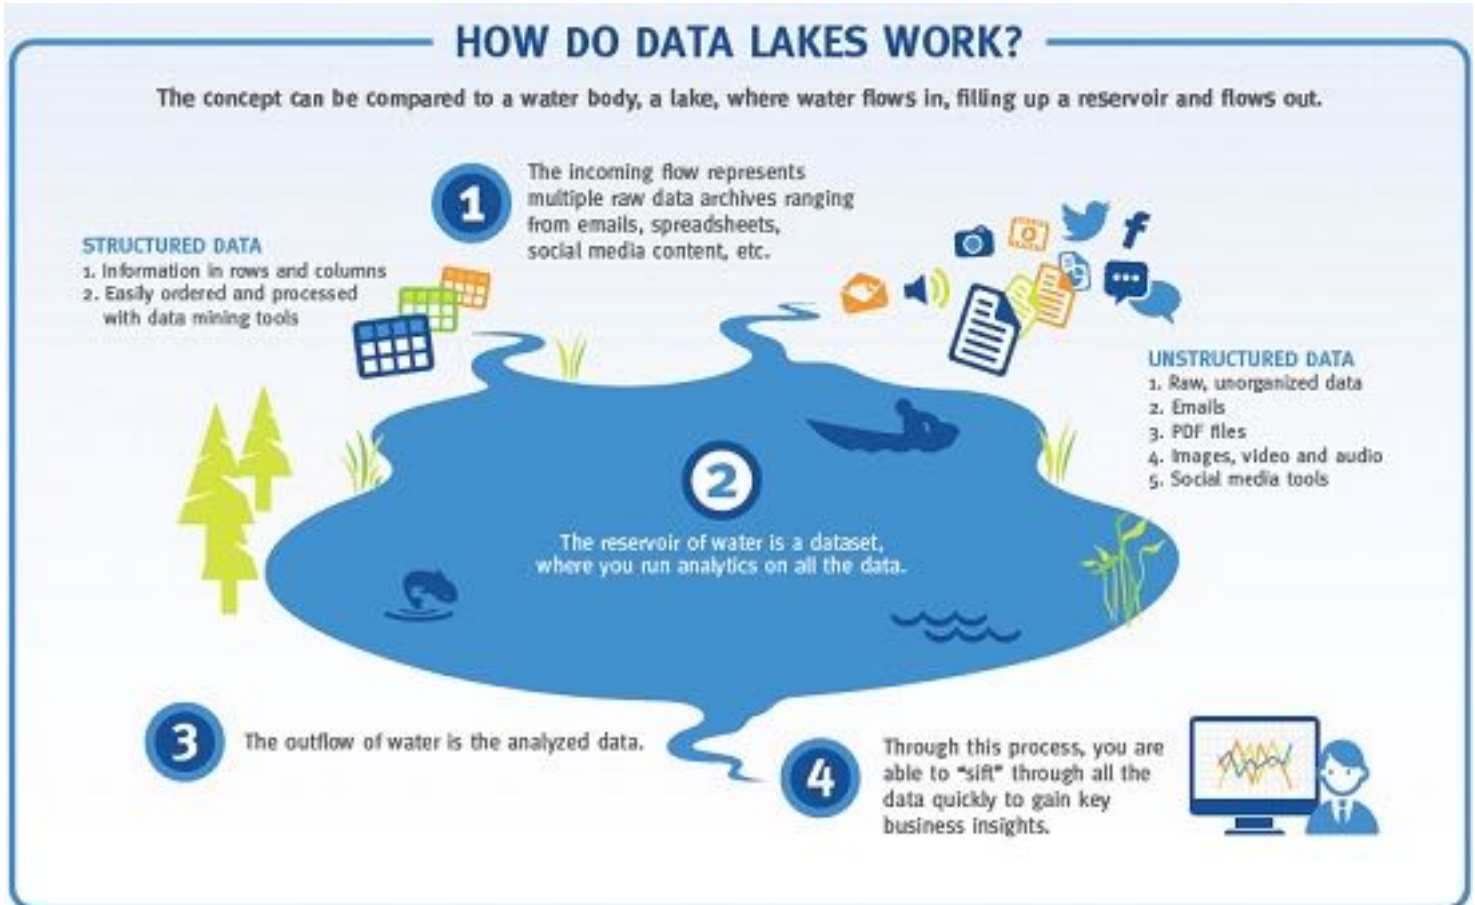
HOW DO DATA LAKES WORK?
The concept can be compared to a water body, a lake, where water flows in, filling up a reservoir and flows out.
1
The incoming flow represents multiple raw data archives ranging from emails, spreadsheets, social media content, etc.
STRUCTURED DATA
1. Information in rows and columns
2. Easily ordered and processed with data mining tools
UNSTRUCTURED DATA
1. Raw, unorganized data
2. Emails
3. PDF files
4. Images, video and audio
5. Social media tools
2
The reservoir of water is a dataset, where you run analytics on all the data.
3
The outflow of water is the analyzed data.
4
Through this process, you are able to "sift" through all the data quickly to gain key business insights.

DATA SCIENCE PART TIME COURSE

# NO-SQL

## SQL

- Traditional rows and columns data
- Strict structure / Primary Keys
- Entire column for each feature
- Industry standard

## NoSQL

- No well defined data structure
- Works better for unstructured data
- Cheaper hardware
- Popular among Startups

## SQL

- MySQL
- Oracle
- Postgres
- SQLite
- SQLServer
- Redshift

## NoSQL

- MongoDB
- CouchDB
- Redis
- Cassandra
- Neo4j
- HBase

# DATA SCIENCE PART TIME COURSE

## JSON - JavaScript Object Notation

- Human readable data with attribute-value pairs.
- What is inside the curly brackets is an object
- In the object we declare variables with ‘attribute’ : ‘value’ pairs

```
var json = {
  "firstName": "John",
  "lastName": "Smith",
  "age": 25,
  "address": {
    "streetAddress": "34 York St",
    "city": "Sydney",
    "state": "NSW",
    "postalCode": "2000"
  },
  "phoneNumbers": [
    {
      "type": "home",
      "number": "02 95999999"
    },
    {
      "type": "office",
      "number": "0431 111 111"
    }
  ],
  "children": [],
  "spouse": null
}
```

- Webservices provide application programming interfaces (APIs) are now usually transferring data via JSON
- Underlying document databases like MongoDB
- Increasingly common data format

# BIG DATA

Data comes from people, technology systems and sensors in the environment

- Organisations web applications
- External web applications (APIs)
- Devices with sensors (Infrastructure, Internet of Things)

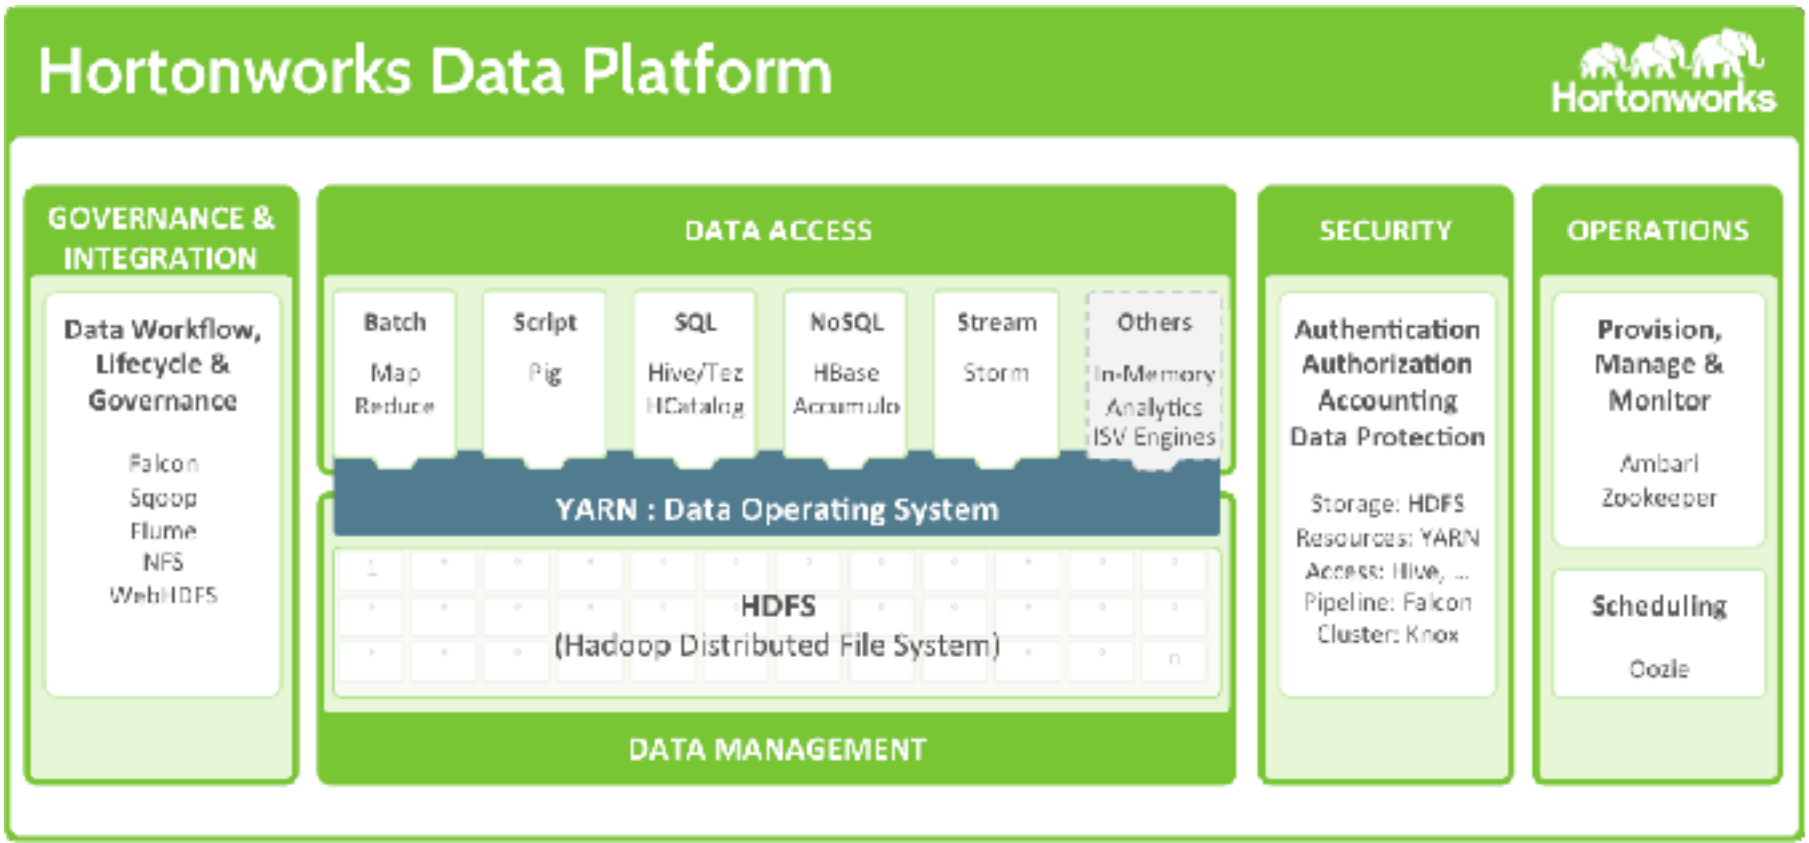
Hortonworks Data Platform
Hortonworks
GOVERNANCE & INTEGRATION
Data Workflow, Lifecycle & Governance
Falcon
Sqoop
Flume
NFS
WebHDFS
DATA ACCESS
Batch
Map Reduce
Script
Pig
SQL
Hive/Tez
HCatalog
NoSQL
HBase
Accumulo
Stream
Storm
Others
In-Memory Analytics
ISV Engines
YARN : Data Operating System
HDFS
(Hadoop Distributed File System)
DATA MANAGEMENT
SECURITY
Authentication
Authorization
Accounting
Data Protection
Storage: HDFS
Resources: YARN
Access: Hive, ...
Pipeline: Falcon
Cluster: Knox
OPERATIONS
Provision, Manage & Monitor
Ambari
Zookeeper
Scheduling
Oozie

cloudera

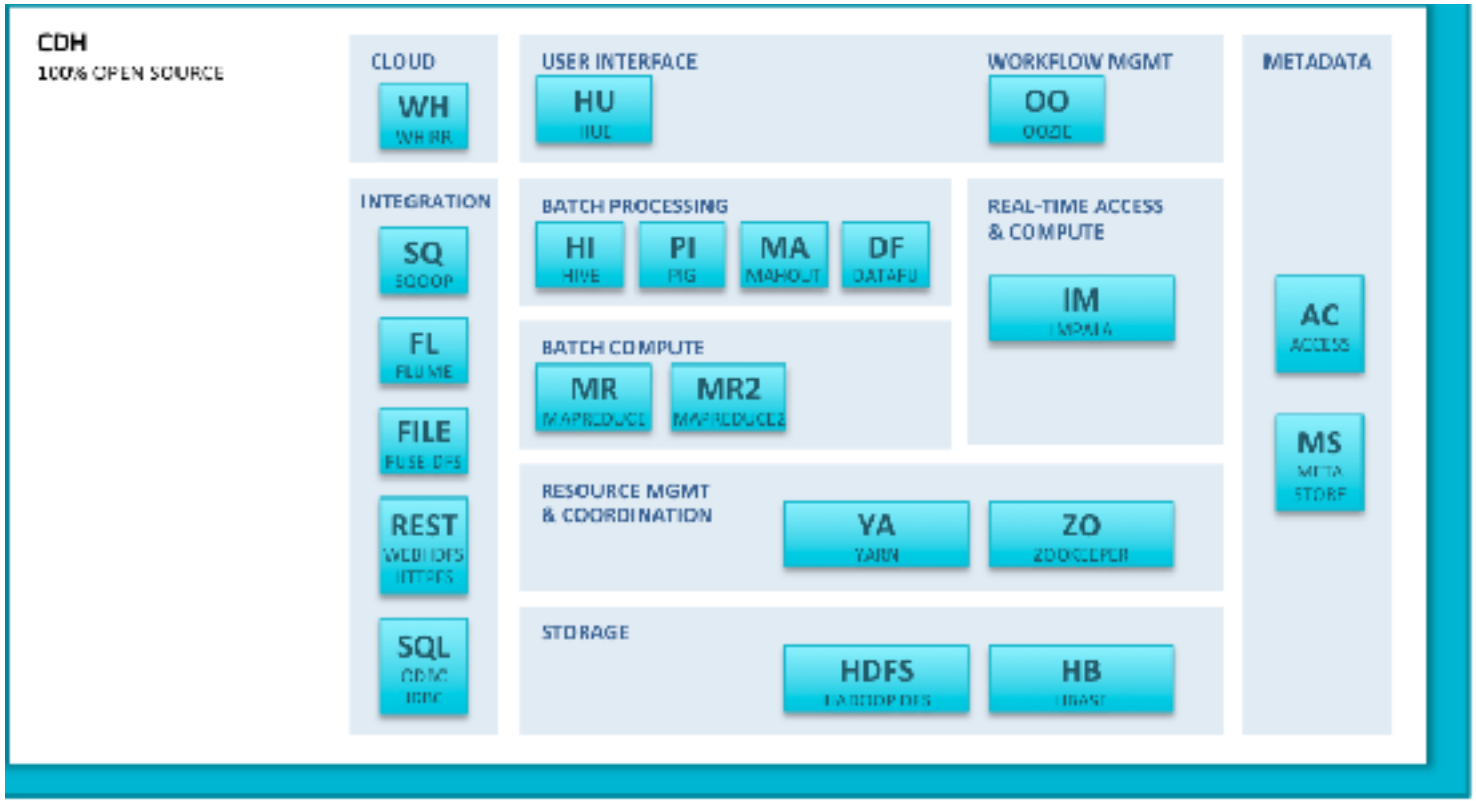
CDH
100% OPEN SOURCE
CLOUD
WH
USER INTERFACE
HU
WORKFLOW MGMT
OO
METADATA
INTEGRATION
SQ
SQOOP
FL
FLUME
FILE
REST
SQL
BATCH PROCESSING
HI
HIVE
PI
PIG
MA
MAHOUT
DF
REAL-TIME ACCESS
& COMPUTE
IM
AC
ACCESS
MS
BATCH COMPUTE
MR
MR2
RESOURCE MGMT
& COORDINATION
YA
YARN
ZO
STORAGE
HDFS
HB

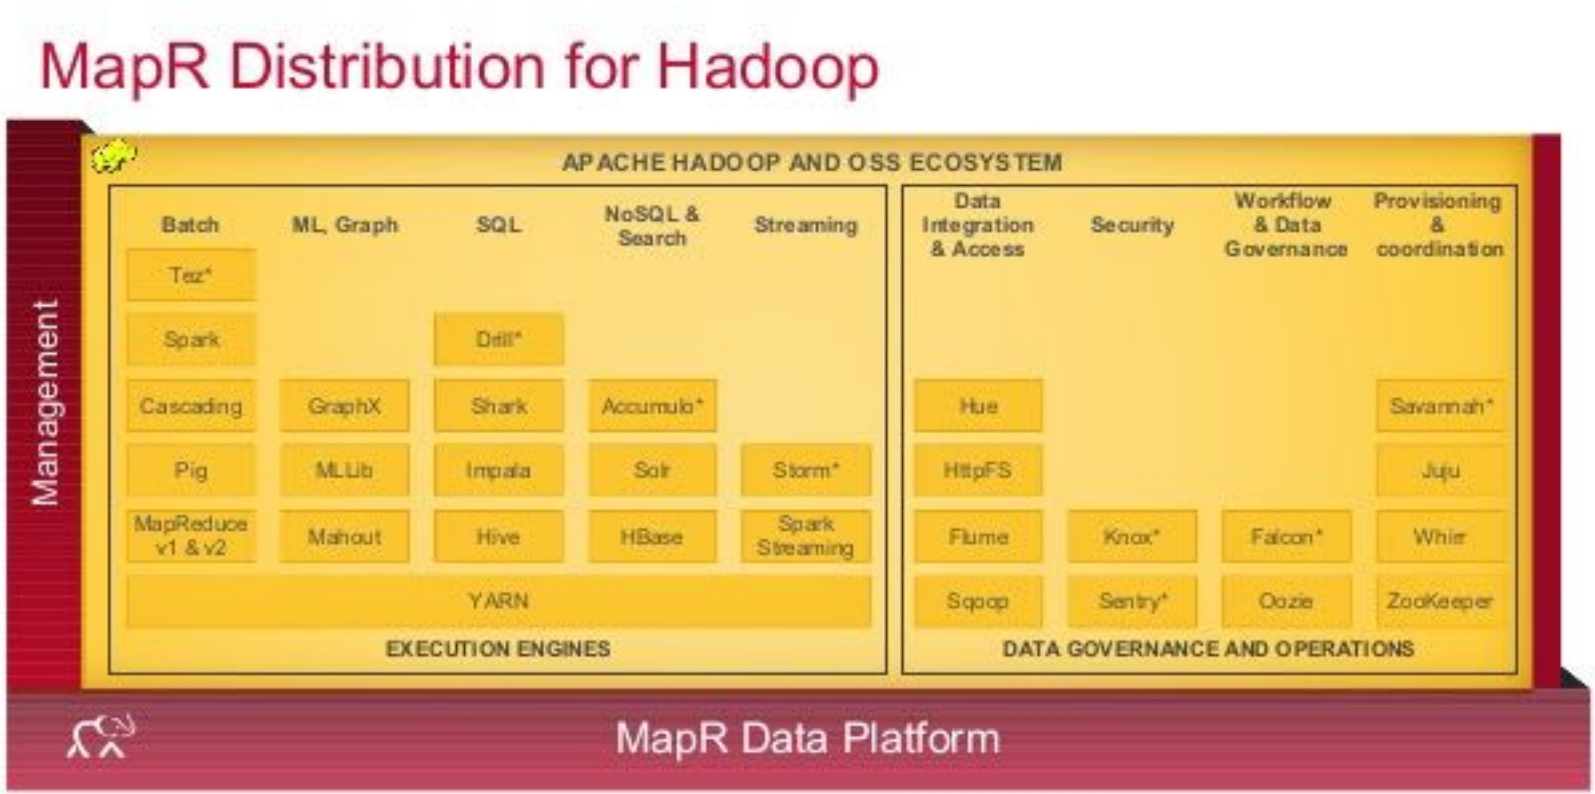
MapR Distribution for Hadoop
APACHE HADOOP AND OSS ECOSYSTEM
Management
Batch
ML, Graph
SQL
NoSQL & Search
Streaming
Tez*
Spark
Drill*
Cascading
GraphX
Shark
Accumulo*
Pig
MLLib
Impala
Solr
Storm*
MapReduce v1 & v2
Mahout
Hive
HBase
Spark Streaming
YARN
EXECUTION ENGINES
Data Integration & Access
Security
Workflow & Data Governance
Provisioning & coordination
Hue
Savannah*
HttpFS
Juju
Flume
Knox*
Falcon*
Whirr
Sqoop
Sentry*
Oozie
ZooKeeper
DATA GOVERNANCE AND OPERATIONS
MapR Data Platform


Avro stores the data definition in JSON format making it easy to read and interpret, the data itself is stored in binary format making it compact and efficient. Avro files include markers that can be used to splitting large data sets into subsets suitable for MapReduce processing.


a new columnar storage format for Hadoop. Parquet can handle complex nested data structures.


Optimized Row Columnar (ORC). A self-describing type-aware columnar file format designed for Hadoop ecosystem workloads. The columnar format lets the reader read, decompress, and process only the columns that are required for the current query.

DATA SCIENCE PART TIME COURSE

# SPARK

Spark is a fast and general processing engine compatible with Hadoop data. It can process data in HDFS, HBase, Cassandra, Hive, and any Hadoop InputFormat. It is designed to perform both batch processing (similar to MapReduce) and new workloads like streaming, interactive queries, and machine learning.

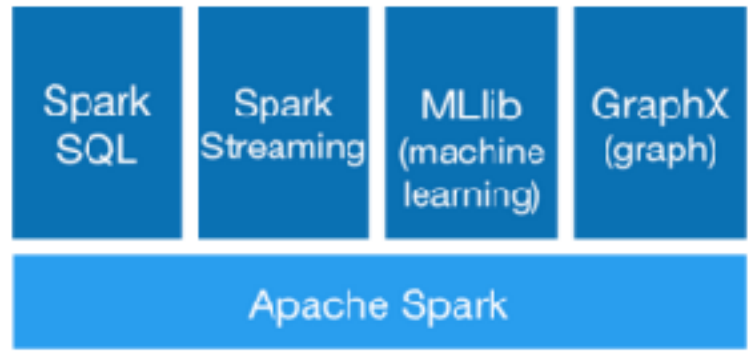


# SPARK – CONTRIBUTORS PER MONTH



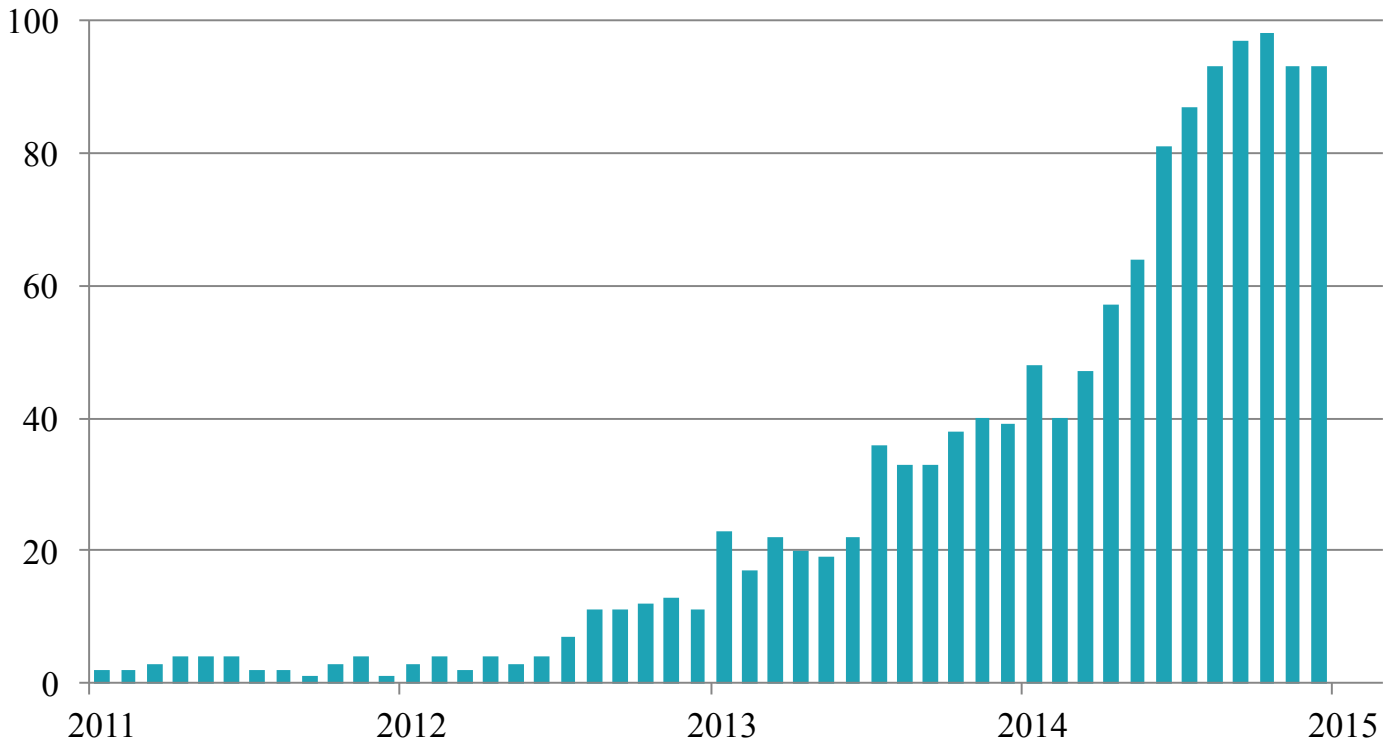

# SPARK – WHAT IS IT?



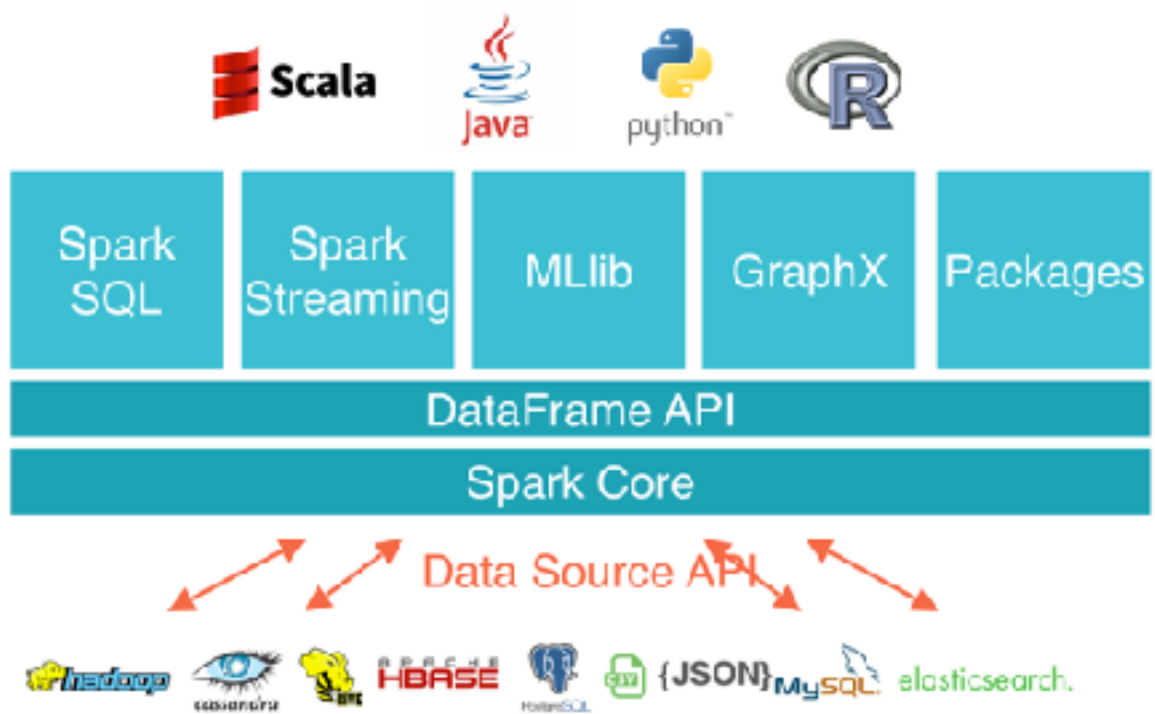

# SPARK – WHAT IS IT?



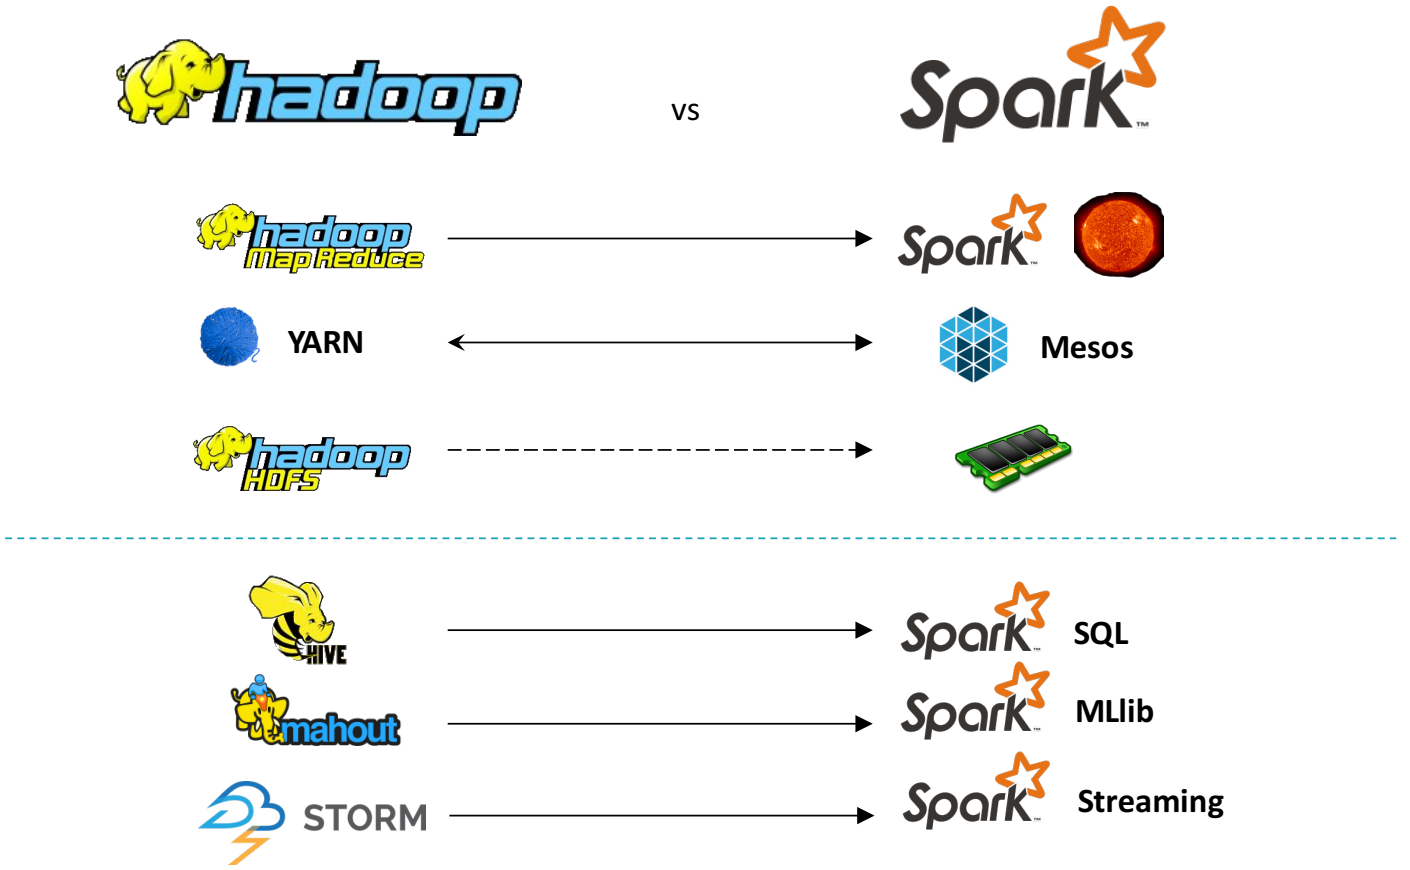

Spark revolves around the concept of a resilient distributed dataset (RDD), which is a fault-tolerant collection of elements that can be operated on in parallel.

There are two ways to create RDDs:

1. Parallelizing an existing collection in your driver program
2. Referencing a dataset in an external storage system, such as a shared filesystem, HDFS, HBase, or any data source offering a Hadoop InputFormat

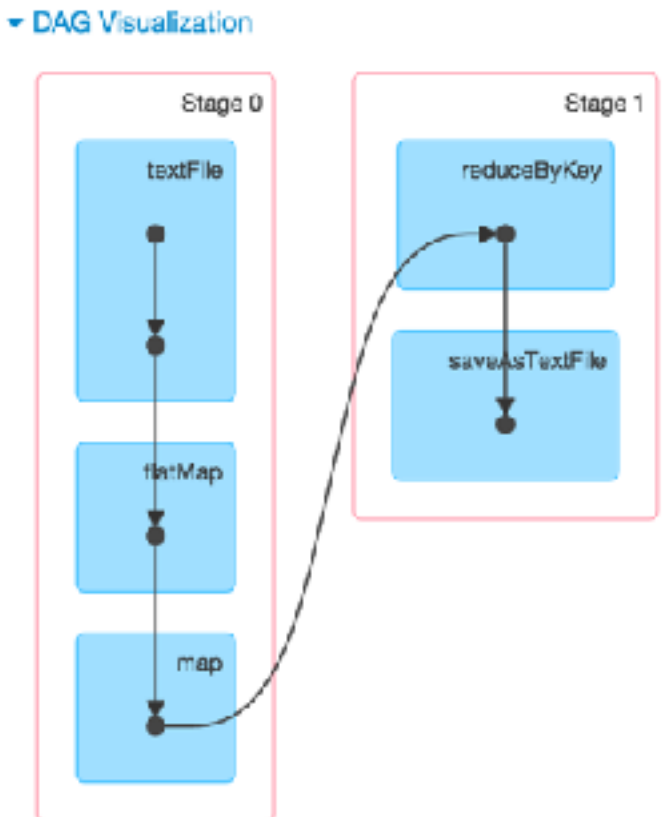
DAG Visualization
Stage 0
textFile
flatMap
map
Stage 1
reduceByKey
saveAsTextFile

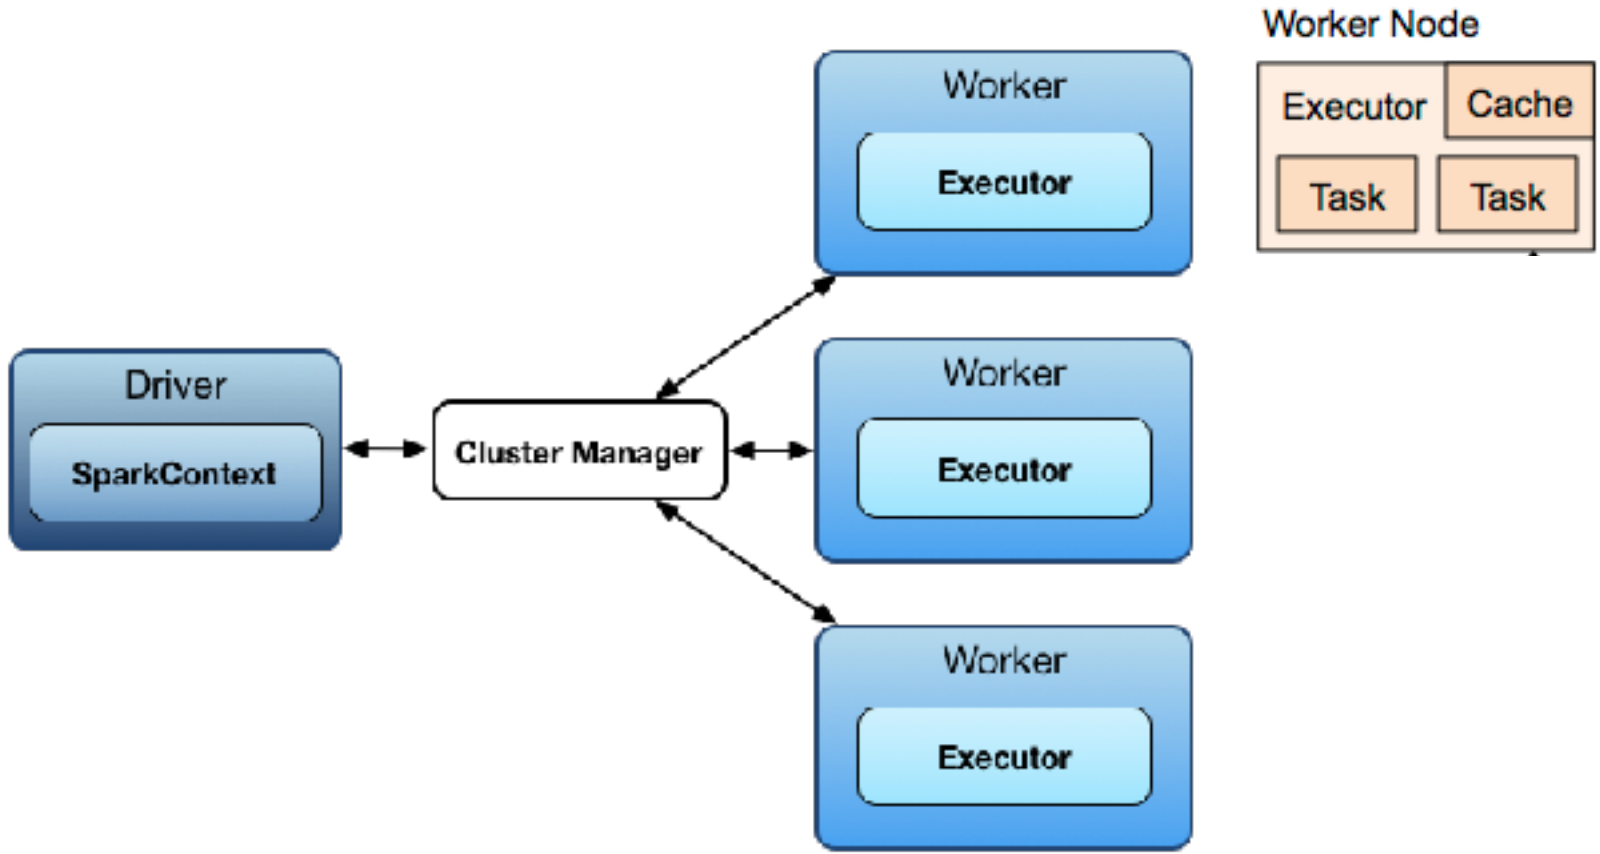
Worker Node
Executor
Cache
Task
Task
Worker
Executor
Driver
SparkContext
Cluster Manager
Worker
Executor
Worker
Executor

DataFrames API was inspired by data frames in R and Pandas in Python. DataFrames integrates with Python, Java, Scala and R.

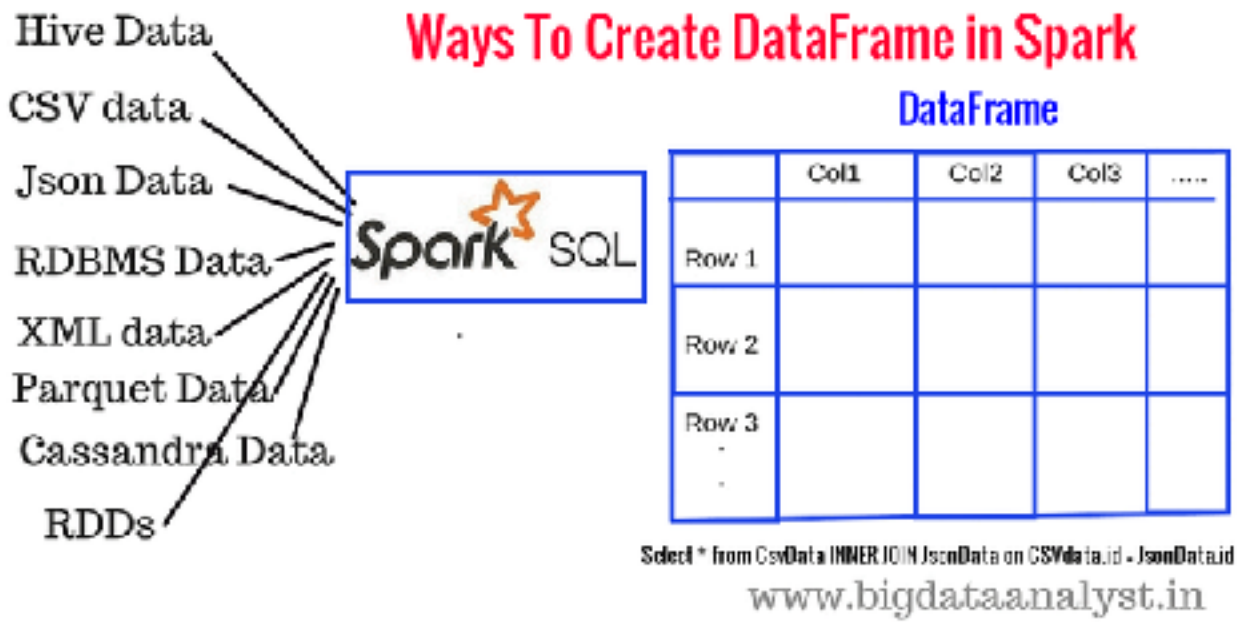


Operations = TRANSFORMATIONS +

ACTIONS

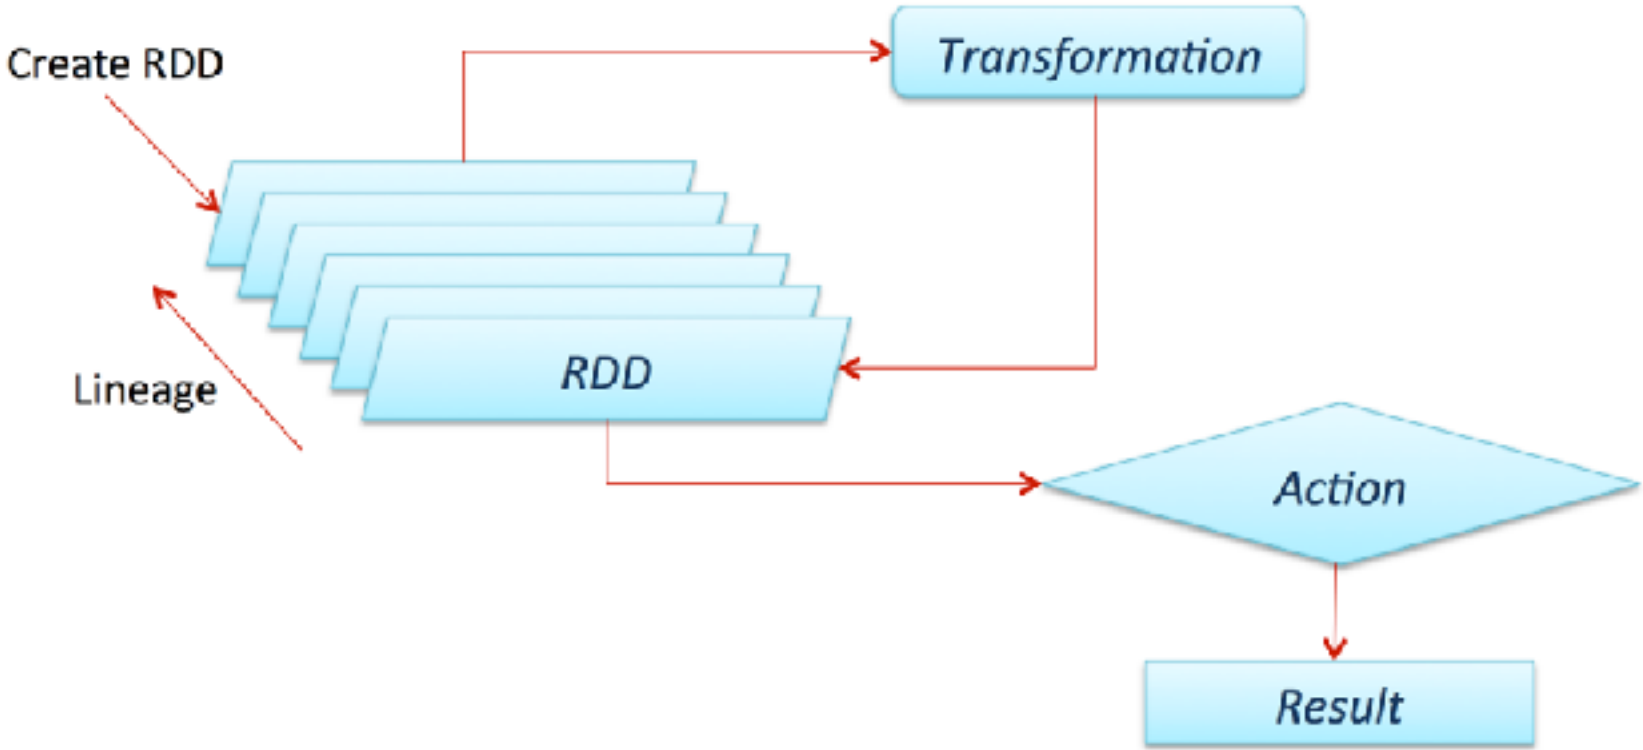
Create RDD
Transformation
RDD
Lineage
Action
Result

 = easy  = medium

## Essential Core & Intermediate Spark Operations


**General**

- map
- filter
- flatMap
- mapPartitions
- mapPartitionsWithIndex
- groupBy
- sortBy

**Math / Statistical**

- sample
- randomSplit

**Set Theory / Relational**

- union
- intersection
- subtract
- distinct
- cartesian
- zip

**Data Structure / I/O**

- keyBy
- zipWithIndex
- zipWithUniqueID
- zipPartitions
- coalesce
- repartition
- repartitionAndSortWithinPartitions
- pipe

---


- reduce
- collect
- aggregate
- fold
- first
- take
- forEach
- top
- treeAggregate
- treeReduce
- forEachPartition
- collectAsMap

- count
- takeSample
- max
- min
- sum
- histogram
- mean
- variance
- stdev
- sampleVariance
- countApprox
- countApproxDistinct

- takeOrdered

- saveAsTextFile
- saveAsSequenceFile
- saveAsObjectFile
- saveAsHadoopDataset
- saveAsHadoopFile
- saveAsNewAPIHadoopDataset
- saveAsNewAPIHadoopFile

Use the programming guides on the Spark Website:

http://spark.apache.org/docs/latest/programming-guide.html

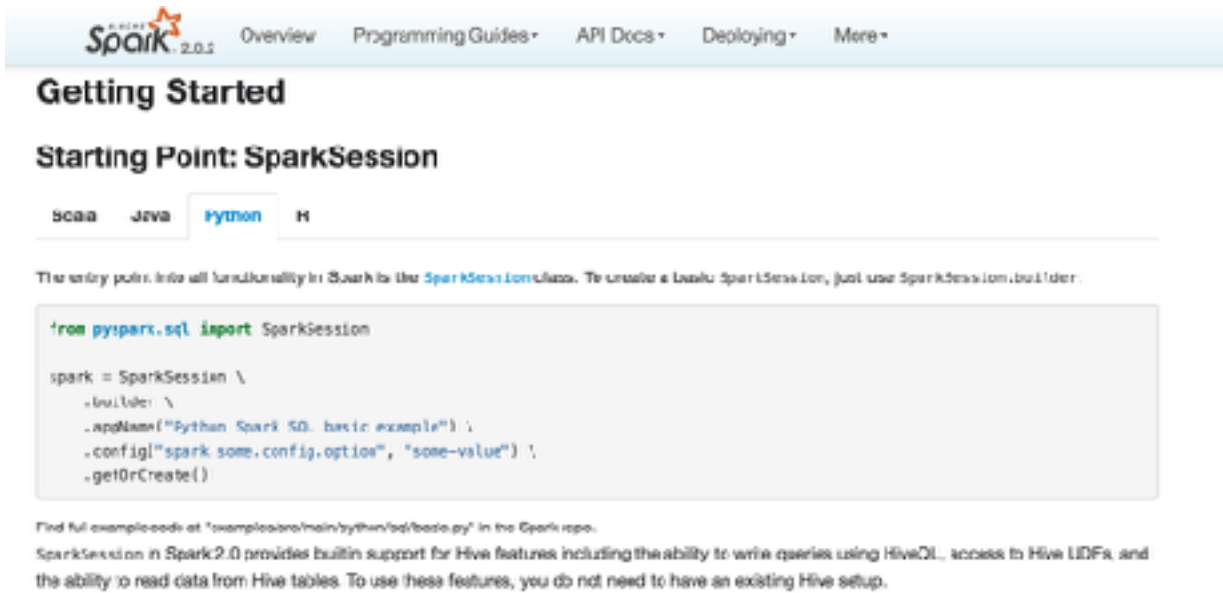


**You can select which programming language API you want to write Spark code with :**

- **Scala**
- **Jave**
- **Python (PySpark)**
- **R (SparkR)**

**Data types**

**Basic statistics**

- summary statistics
- correlations
- stratified sampling
- hypothesis testing
- streaming significance testing
- random data generation

**Classification and regression**

- linear models (SVMs, logistic regression, linear regression)
- naive Bayes
- decision trees
- ensembles of trees (Random Forests and Gradient-Boosted Trees)
- isotonic regression

**Collaborative filtering**

- alternating least squares (ALS)

**Clustering**

- k-means
- Gaussian mixture
- power iteration clustering (PIC)
- latent Dirichlet allocation (LDA)
- bisecting k-means
- streaming k-means

**Dimensionality reduction**

- singular value decomposition (SVD)
- principal component analysis (PCA)

**Feature extraction and transformation**

**Frequent pattern mining**

- FP-growth
- association rules
- PrefixSpan

**Evaluation metrics**

**PMML model export**

**Optimization (developer)**

# COURSE FEEDBACK

# LAB

- Log into Databricks community edition https://databricks.com/try-databricks
- https://docs.databricks.com/_static/notebooks/gentle-introduction-to-apache-spark.html

# In databricks import the notebook from the below URL



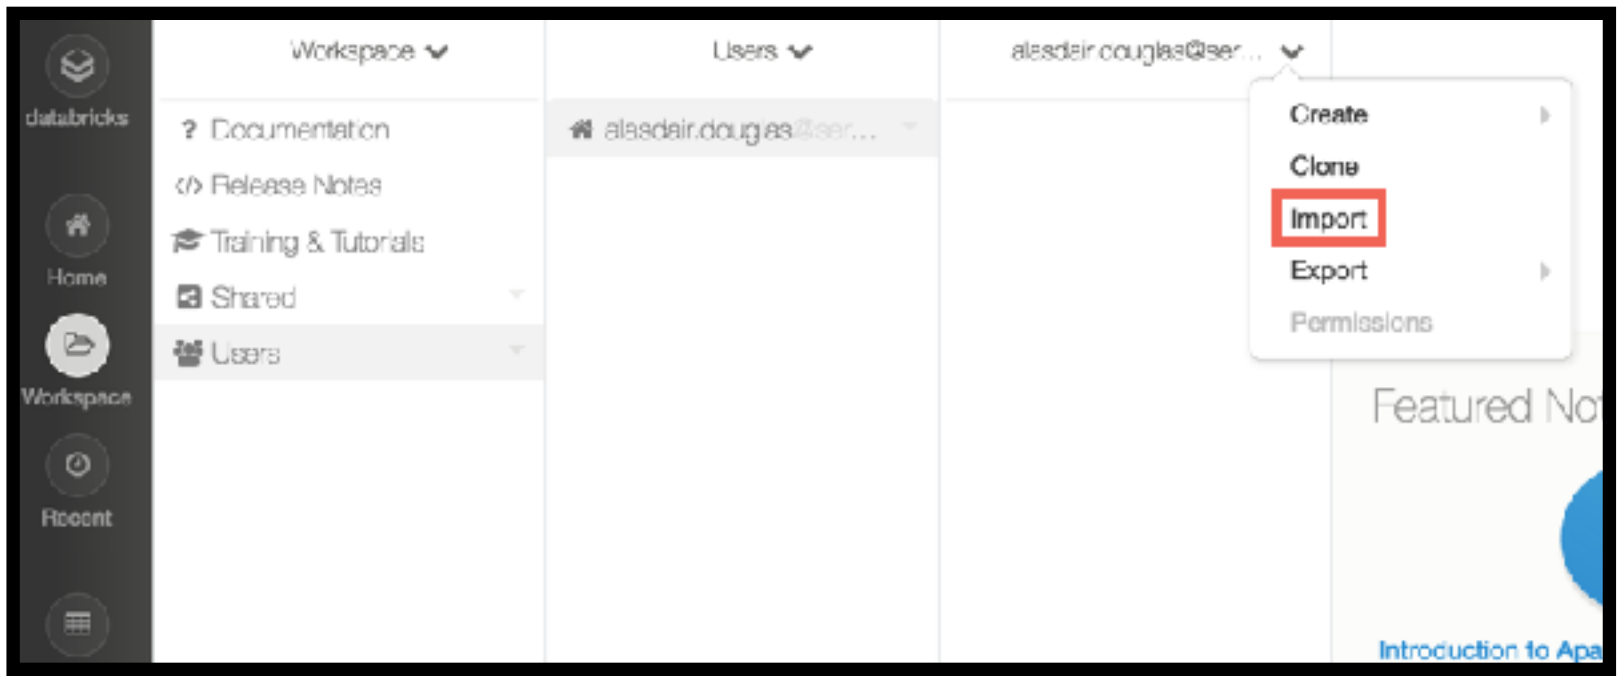


https://docs.databricks.com/_static/notebooks/gentle-introduction-to-apache-spark.html

# SYNCHING YOUR FORK WITH THE COURSE REPO



1. re-name your labs with lab_name.<yourname>.ipynb  (to prevent a conflict)
2. cd <path to the root of your SYD_DAT_6 local repo>
3. commit your changes ahead of sync
   - git status
   - git add .
   - git commit -m "descriptive label for the commit"
   - git status
4. download new material from official course repo (upstream) and merge it
   - git checkout master  (ensures you are in the master branch)
   - git fetch upstream
   - git merge upstream/master

# HOMEWORK

**Homework**

- Read Natural Language Processing website – http://www.nltk.org/ (5 mins)
- Read and be able to explain one use case of the Alchemy API (10 mins)
- Download and install NLTK for Python (10mins)

**Reading**

- For those that want a good foundation in data stores and architecture, the redbook 5th edition is a good reference http://www.redbook.io/